

Programming Assignment 2 : Group 3

Members : Judah, Ramon, Siddhesh, Yaksh

Our Results :

1. Simulation Component

(a). We have choose to implement SACK and using an initial CW = 4 for this assignment.

Error percentage = 2.5%

Wireless Link = 8 Mbps

The reasons have been explained in 1.(c) and 1.(d).


End to end throughput with and without the TCP option is as follows (for a simulation period of 100 seconds and actually transmission time of 99 seconds):

Case :	Throughput (Mbps)
With SACK disabled and CW = 1 (No TCP-option implemented)	1.2390
With SACK disabled and CW = 4	1.4070
With SACK enabled and CW = 1	1.7803
With SACK enabled and CW = 4 (Both TCP-option implemented)	1.9202


(b). Instantaneous Throughput (averaged over 25 ms)

Chart 1 and Chart 2 give the instantaneous throughput results. We have simulated for 100 seconds (actual transmission time is 99 seconds), but we have intentionally selected a smaller period to plot the instantaneous throughput charts below as using all 99 seconds of instantaneous throughput results led to a non-understandable, highly crowded graph. Also, we are attaching the complete instantaneous throughput calculation file in our .zip folder for your reference. These instantaneous throughput plots were changed from 100 ms to 25ms for better interpretation and to show the saw tooth nature of TCP. For the same reason, we are also plotting the CW as a better representation of the saw tooth nature we observed.


Regular TCP Throughput




Increased Initial CW Throughput




SACK Implemented Throughput




Both TCP Options Enabled Throughput




Regular TCP Congestion Window



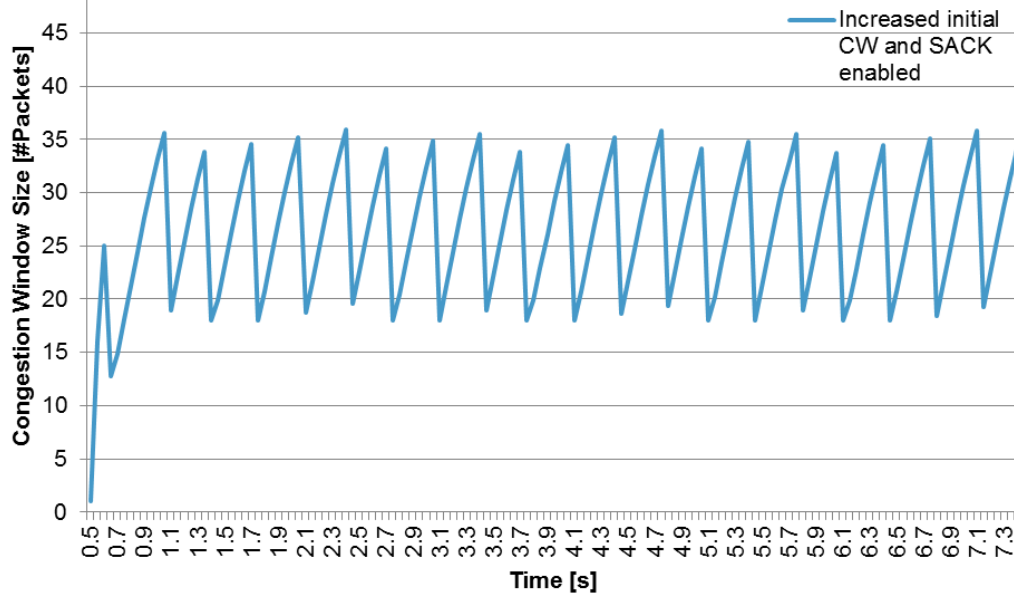
Increased Initial CW Congestion Window



SACK Implemented Congestion Window



Both TCP Options Enabled Congestion Window



(c). Give reasons for the choice of 2 options and the corresponding network conditions.

Choice of the 2 options of TCP :

SACK

Without SACK enabled, TCP requires one RTT to recover from each packet loss and this is a big penalty to pay, especially in wireless communication wherein loss is substantially more than wired, often leading to multiple packet loss (due to random loss, burst loss etc). Here, it would take $n \times \text{RTT}$ to recover from n packets lost - an inefficient mechanism for wireless communication.

With SACK enabled, it allows the transmitter to recover from more than one packet loss in one RTT. In wireless communication, we often see multiple packets losses (especially burst packet losses) and by enabling SACK we are able to detect and recover from these multi-packet losses faster, therefore increasing the performance of TCP.

CW = 4

TCP connection is vulnerable in the beginning of slow start phase ($CW = 1$ or 2). A packet loss during this time window will require TCP to wait for a timeout as there are not enough packets transmitted to cause dupacks. In wireline connection, this is not a problem as losses do not normally occur and unless there is congestion in the network, in which case, it is right for the connection to start out slowly and continue with a low CW until congestion clears out. In a wireless connection, however, a loss can occur for reasons other than network congestions (e.g. random, handoff). It is therefore not uncommon for a packet loss to occur in the beginning of slow start. If such a case occurs, TCP will have to wait for a timeout in order to recover from the loss. The time in which TCP waits for a timeout could have been used to build up the congestion window. By changing the initial congestion window into 4, we eliminate this problem. With initial congestion window set to 4, when a random packet loss occur during the beginning of slow start, there will be other packets transmitted to trigger dupacks and TCP will be able to recover through fast retransmit. In this way, TCP can continue increasing its congestion window and reach its optimal size. This performance improvement at the beginning of the connection is especially important for wireless connections in which connections are usually short-lived.

Network Conditions

We have selected the following network conditions :

Error percentage = 2.5%

Wireless Link = 8 Mbps

Delay of the link = 5 ms

The reason for selecting an error percentage of 2.5% was that it was within the practical range of losses that are experienced practically in wireless connection. By selecting a loss rate lower, we would not have been able to show the effectiveness of the two options that we have enabled. By having too high of loss will lead to a substantial increase in the effectiveness of the TCP options over the no TCP options, however, the raw throughput of the system will be too low for a communication protocol (we need to show that not only is performance increase substantial, but the utilization of the link is also not too low). Hence we have chosen the intermediate value of 2.5%.

For the bandwidth of the wireless link, we have chosen 8Mbps considering the following factors. First, the wired link is 20Mbps and in any practical situation, we always have the wired link to have a sufficiently greater bandwidth than the wireless bandwidth. Hence our wireless link would have to be substantially less than the wired 20Mbps. Next we compared the performance of the wireless link (throughput comparisons) for different BW choices and observe that for the loss = 2.5% , the throughput of the link does not substantially increase with an increase in bandwidth after a certain BW (here, it was 8Mbps). In other words, we have considered the spectral efficiency to be a factor in selecting the BW of the wireless link. For example, we feel it would be a highly inefficient utilization of the spectrum (a precious resource) to double the bandwidth (to 16Mbps) for a marginal increase in throughput of 25%. We found that from this justification that 8Mbps was the optimal utilization of the bandwidth for the assumed network conditions.

As for the delay of the wireless link channel, we looked at practical scenarios for the average delay of a link, such as the RTT of ping requests over WiFi in different locations and observed that an average delay of less than 1ms was observed in the links. In order to simulate an unfavorable link that leads to a bottle neck at Node 2 due to delays , we wished to select a value that was higher than 2mS. We found from literature that delays up to the range of 8 to 10 ms were not unheard of for a single hop and therefore we took an intermediate value of 5ms to avoid being on extreme ends of the observed values.

Reasoning of the selection of the two options implemented over other options

(we initially interpreted 1(d) to answer this question; we have kept this answer as an additional question):

We feel our choice of options is appropriate as the aim was to choose two options that would have the most significant positive impact on the link performance in these network conditions.

The other options, while all have substantial merit, were not in the top two options to be implemented for this assignment for the following reasons :

Limited transmit : By enabling SACK option, we will be notifying the TX about multiple packet loss in one single ACK packet. This allows the TX to have multiple opportunities to be able to transmit the same packet again, thereby reducing the chances of DUPACKs to mainly those cases where the packet lost is at the beginning of the CW. In the majority of the other cases, due to SACK, we would have captured and prevented 3 Dupacks from occurring and therefore the effectiveness of this option would have not have been as substantial as using another option. Hence we choose SACK over Limited Transmit option.

Timestamp option : Using the timestamp option for every packet in order to calculate a more accurate RTT leads a significant overhead and processing power. Enabling the other options will have a more impact on throughput than having accurate RTT information. Also, using accurate RTT of last packer might not always bring about positive results as the channel condition at a given point of time is not necessarily a function of the past condition, i.e previous packet's RTT may not have any relation to the next packet's channel condition (although the network infrastructure will be nearly the same for consecutive packets, we wish to say that the channel conditions may not be similar for consecutive packets).

Disabling header compression : We wish to consider the network conditions as an important factor and that a lossy link should always perform worse than an ideal link (no loss). If we enable the option of disabling header compression, in situations where the link quality does improve to close to ideal conditions, this option will actually limit the throughput. In fact, in an ideal no loss channel, this option will reduce throughput (as compared to header compression present in regular TCP) due to the increased overhead associated with every packet. Therefore, although each option has its merits, we feel these other options will not be as effectual as the two that were chosen.

1.(d)

From our results we can see that the throughput of the link has increased due to the implementation of the TCP options. For each of the option, we notice that :

2. SACK : For SACK being enabled, we see that the throughput has substantially increased over no SACK. This can be attributed to less congestions occurring due to the selective ACKs reducing 3 DUPACKs from causing congestion, and thereby allowing the CW to increase. Maintaining a high CW and sustaining it for a longer time period allows the substantial increase in throughput.
3. CW = 4. We observe an increase in throughput here as well and notice that it is not as large as SACK. The throughput increase is due to the larger CW which prevents timeouts from occurring when the CW is less than 4. As a result, traffic is continuously flowing in to the network, even after a timeout, unlike when CW =1, where one packet loss creates a timeout, due to lack of 3 Dupacks. Also, we notice that the throughput increase is not as high as SACK option enabled as the CW = 4 option takes effect after congestion, in the slow start mode only, whereas SACK's effect is present throughout all three modes of the TCP connection, therefor it leads a higher increase in throughput.

WIRESHARK ANSWERS

2a.

local host: 192.168.1.187

remote host: 74.125.212.177

filter: http.request.uri


uri:

/videoplayback?sparams=id%2CexpilYNUluQ0pPVXMLAAAAaEpreV8zbkR0cElzCw

(TCP Stream FLOW)

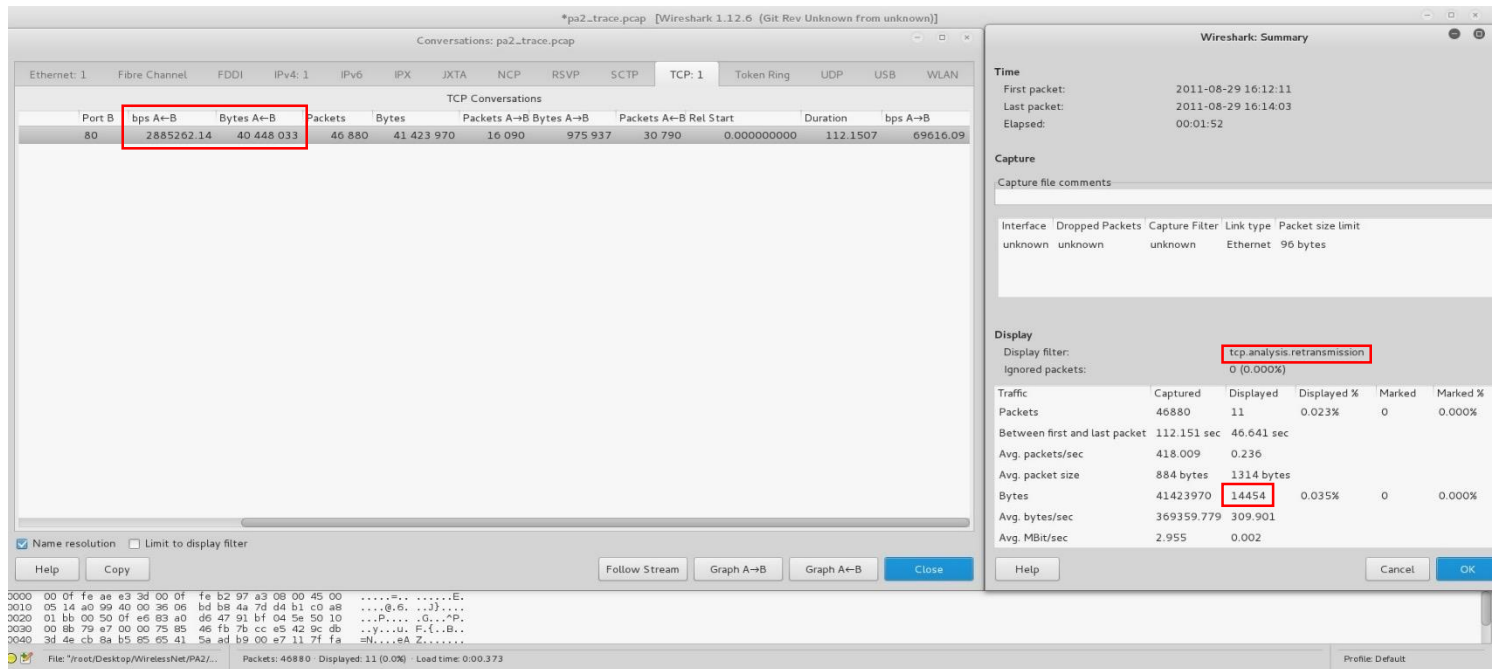
2b.

RTT = Filter:tcp.analysis.ack_rtt

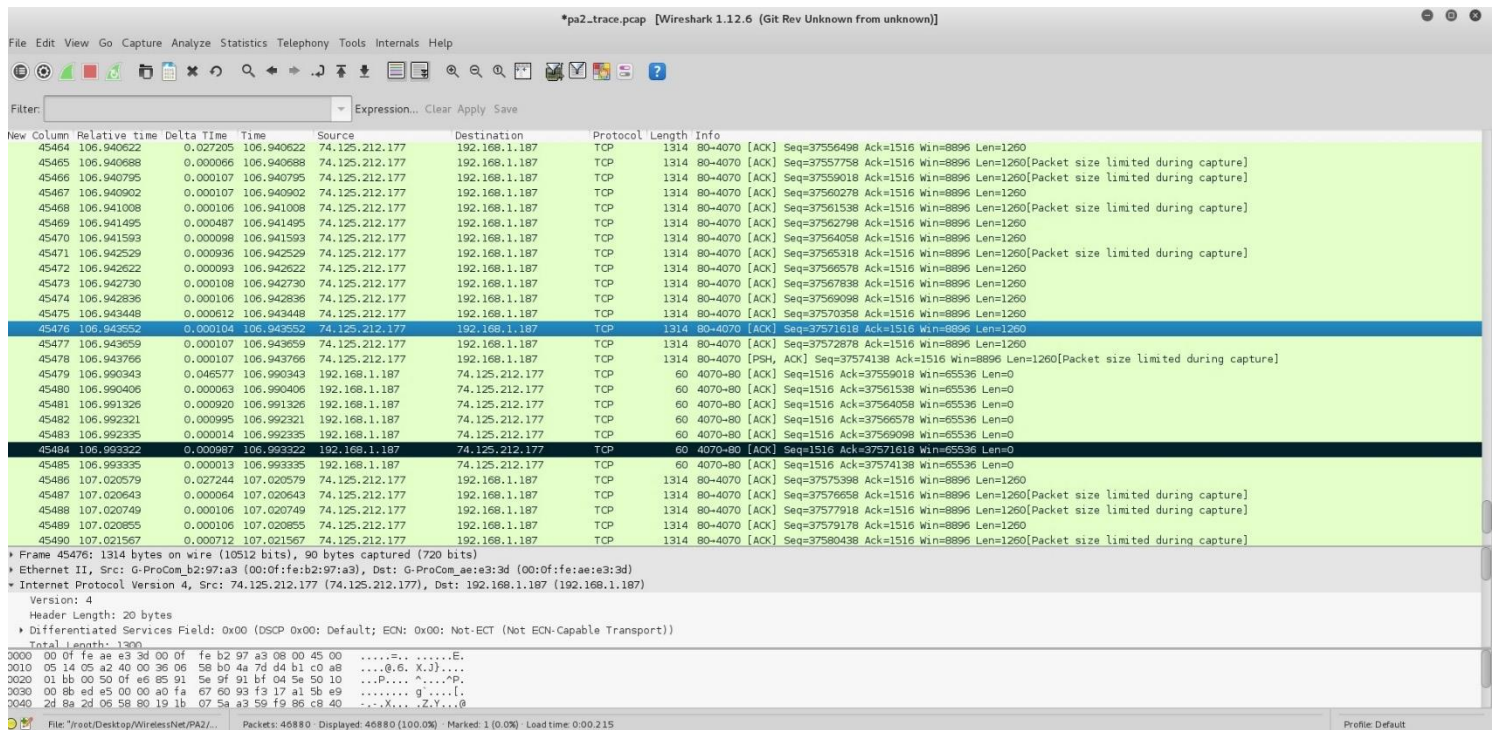


For actual downlink throughput=

$$\left(\frac{((\text{Total \#bytes} - \text{\#retx bytes}) * 8)}{(1,000,000 * \text{Total Duration})} \right) \text{ Mbps}$$
$$= \left(\frac{((40448033 - 14454) * 8)}{(1,000,000 * 112.057)} \right)$$
$$= 2.88423 \text{ Mbps}$$



To confirm that the Avg. Throughput value isn't off by a lot, we can calculate instantaneous throughput of one packet with respect to its sequence and ACK Number. In the figure below, packet with sequence number 37571618 is ACKed at 106.993322, so Throughput= 2.8092 Mbps, which is not that different from the Avg. Throughput.

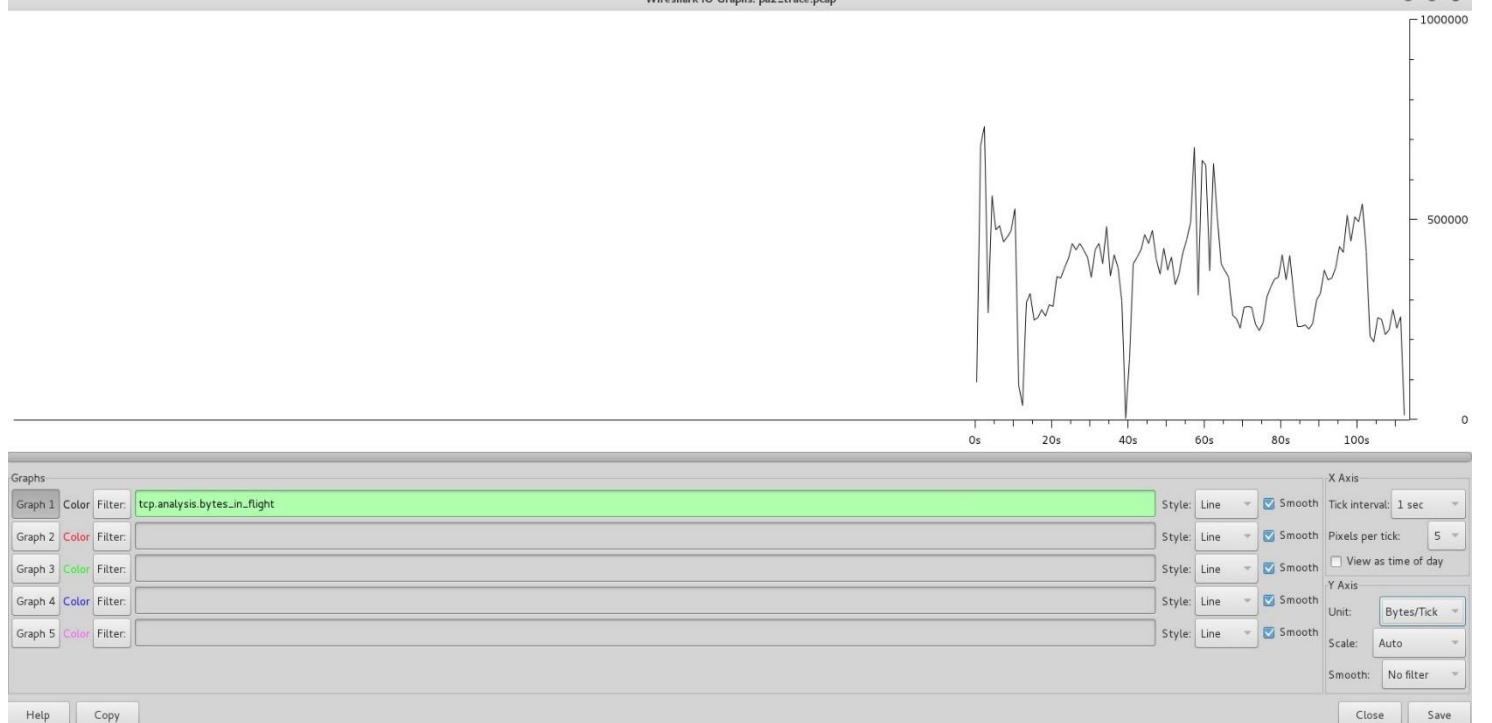
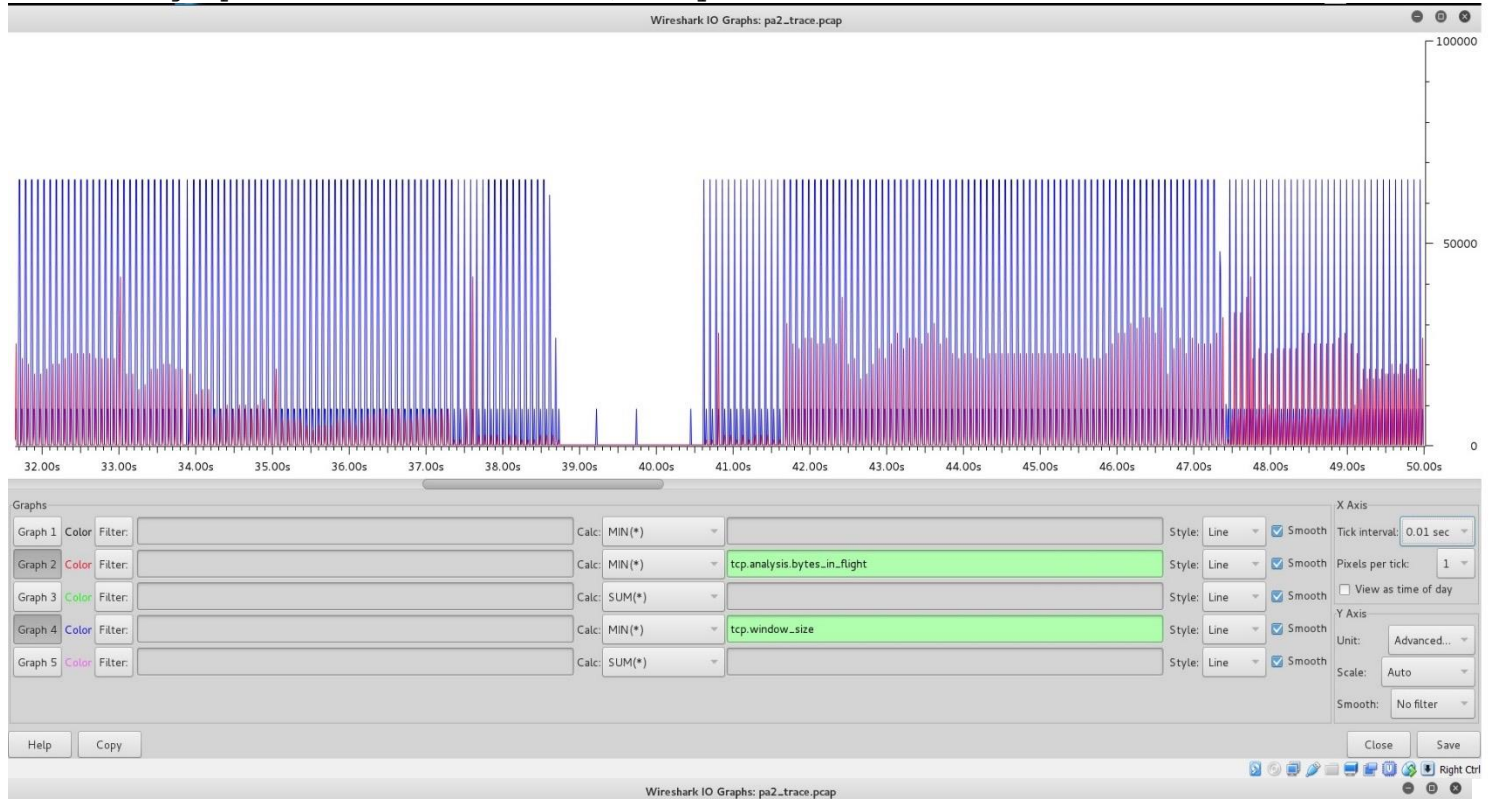


Congestion Window:



Filter: `tcp.analysis.bytes_in_flight`

We can't 'get' that value directly from the capture file, as it is not advertized. `bytes_in_flight` are un-ACKed bytes for Wireshark. So, this gives us a good estimate of the congestion window progression. Also, if `bytes_in_flight` is below the advertized window size, it can also mean that the receiver ACKed before the receive window was full.

Both the graphs below are the same, only with different Tick Interval.



Another Graph, time-sequence number graph, can give a better idea about congestion and how fast is the recovery. Both Plots below are the same with the second plot being a zoomed in version of the first one around time=40.0 Seconds where Zero Window Occurs.



2c.

Zero Window

filter: tcp.analysis.zero_window

From Receivers point of view (192.168.1.187,4070,74.125.212.177,80)

Receivers Side:

Through TCP ACK frames, the client informs the server of how much room is in this buffer. If the TCP Window Size goes down to 0, the client will not be able to receive any more data until it processes and opens the buffer up again. In this case, Protocol Expert will alert a "Zero Window" in Expert View. The receiver keeps doing this till it can process more data.

Possible Reason for Zero Window:

It could be that the machine is running too many processes at that moment, and its processor is maxed. Or it could be that there is an error in the TCP receiver, like a Windows registry misconfiguration.

Senders Side:

The sending station is trying to see if the remote peer is dead(using Keep Alives), if the connection is still open and in use, or may just need to keep the connection open instead of suffering another handshake overhead. If the target does not respond, the sender may send several Keep Alives before finally sending a TCP reset to kill the socket. This is a good thing, since we don't want open/unused TCP connections staying open and hogging resources forever.



pa2_trace.pcap [Wireshark 1.12.6 (Git Rev Unknown from unknown)]

Filter: tcp.analysis.zero_window

No.	Time	Source	Destination	Protocol	Length	Info
16951	38.702778	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=13997018 Win=16396 Len=0
16952	38.703757	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=13997038 Win=13876 Len=0
16953	38.703774	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14002058 Win=11356 Len=0
16954	38.703789	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14004578 Win=8896 Len=0
16955	38.704759	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14007098 Win=6316 Len=0
16956	38.704776	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14009618 Win=3796 Len=0
16957	38.704792	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14012138 Win=1276 Len=0
16958	38.730134	192.168.1.187	74.125.212.177	TCP	1314	80->4070 [ACK] Seq=14012138 Ack=1516 Win=8896 Len=1260[Packet size limited during capture]
16959	38.898722	192.168.1.187	74.125.212.177	TCP	60	4070->80 [ACK] Seq=1516 Ack=14013398 Win=16 Len=0
16960	39.228137	74.125.212.177	192.168.1.187	TCP	70	[TCP Window Full] [TCP segment of a reassembled PDU]
16961	39.401586	192.168.1.187	74.125.212.177	TCP	60	[TCP ZeroWindow] 4070->80 [ACK] Seq=1516 Ack=14013414 Win=0 Len=0
16962	39.746144	74.125.212.177	192.168.1.187	TCP	54	[TCP Keep Alive] 80->4070 [ACK] Seq=14013413 Ack=1516 Win=8896 Len=0
16963	39.795495	192.168.1.187	74.125.212.177	TCP	60	[TCP ZeroWindow] 4070->80 [ACK] Seq=1516 Ack=14013414 Win=0 Len=0
16964	40.455553	74.125.212.177	192.168.1.187	TCP	54	[TCP Keep Alive] 80->4070 [ACK] Seq=14013413 Ack=1516 Win=8896 Len=0
16965	40.504313	192.168.1.187	74.125.212.177	TCP	60	[TCP ZeroWindow] 4070->80 [ACK] Seq=1516 Ack=14013414 Win=0 Len=0
16966	40.625281	192.168.1.187	74.125.212.177	TCP	60	[TCP Window Update] 4070->80 [ACK] Seq=1516 Ack=14013414 Win=5536 Len=0
16967	40.695419	74.125.212.177	192.168.1.187	TCP	1314	80->4070 [ACK] Seq=14013414 Ack=1516 Win=8896 Len=1260
16968	40.695490	74.125.212.177	192.168.1.187	TCP	1314	80->4070 [ACK] Seq=14014674 Ack=1516 Win=8896 Len=1260
16969	40.695594	74.125.212.177	192.168.1.187	TCP	1314	80->4070 [ACK] Seq=14015934 Ack=1516 Win=8896 Len=1260
16970	40.695701	74.125.212.177	192.168.1.187	TCP	1314	80->4070 [ACK] Seq=14017194 Ack=1516 Win=8896 Len=1260
16971	40.695809	74.125.212.177	192.168.1.187	TCP	1314	80->4070 [ACK] Seq=14018454 Ack=1516 Win=8896 Len=1260[Packet size limited during capture]

Frame 16951: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface 0

Ethernet II, Src: G-ProcCom-ae3:3d (00:0f:fa:ae:3:3d), Dst: G-ProcCom-b2:97:a3 (00:0f:fa:b2:97:a3)

Internet Protocol Version 4, Src: 192.168.1.187 (192.168.1.187), Dst: 74.125.212.177 (74.125.212.177)

Transmission Control Protocol, Src Port: 4070 (4070), Dst Port: 80 (80), Seq: 1516, Ack: 13997018, Len: 0

Source Port: 4070 (4070)

Destination Port: 80 (80)

[Stream index: 0]

[TCP Segment Len: 0]

Sequence number: 1516 (relative sequence number)

Acknowledgment number: 13997018 (relative ack number)

Header Length: 20 bytes

... 0000 0001 0000 = Flags: 0x010 (ACK)

0000 00 0f fe b2 97 a3 00 0f fe ae a3 3d 08 00 45 00=...E.
0010 00 28 b3 f6 40 00 80 65 47 c0 a8 01 bb 4a 7d (...@...e@...J)
0020 d4 b1 0f e0 50 91 bf 04 5e 94 29 a6 57 50 10 (......^...)WP.
0030 20 06 d6 d7 60 00 00 00 00 00 00 00 00 00 00 00 ..b.....

File: /root/Desktop/WirelessNet/PA2... Packets: 46880 Displayed: 46880 (100.0%) Load time: 0:00:197

Profile: Default

TCPTRACE ANSWERS

2a.

local host: 192.168.1.187

remote host: 74.125.212.177

Command:

```
tcptrace -e pa2_trace.pcap
```

(for uri extraction, this extracts content of tcp stream into a file which contains the uri)

```
uri: /videoplayback?sparams=id%2CexpilYNUluQ0pPVXMLAAAAaEpreV8zbkR0cElzCw
```

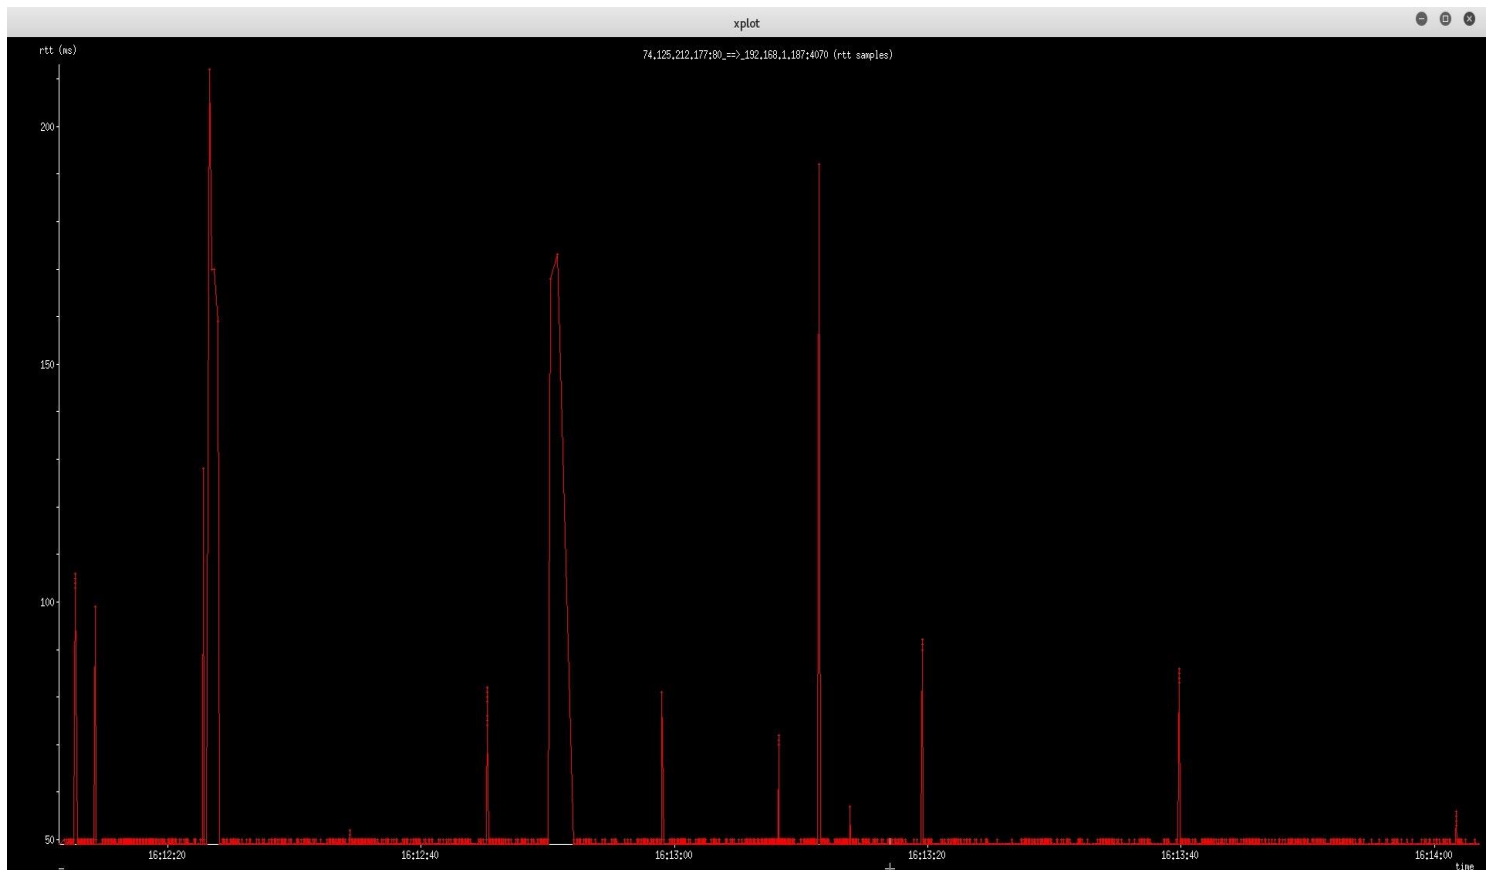
2b.

RTT

Commands:

```
tcptrace -G pa2_trace.pcap (Creates all the graphs)
```

```
xplot.org b2a_rtt.xpl (plots rtt for b->a)
```



Downlink Throughput:

Command:

```
tcptrace -r -l pa2_trace.pcap
```

46880 packets seen, 46880 TCP packets traced

elapsed wallclock time: 0:00:00.052966, 885096 pkts/sec analyzed

trace file elapsed time: 0:01:52.150733

TCP connection info:

1 TCP connection traced:

TCP connection 1:

host a: 192.168.1.187:4070

host b: 74.125.212.177:80

complete conn: yes

first packet: Mon Aug 29 16:12:11.356700 2011

last packet: Mon Aug 29 16:14:03.507433 2011

elapsed time: 0:01:52.150733

total packets: 46880

filename: pa2_trace.pcap

a->b:	b->a:		
total packets:	16090	total packets:	30790
ack pkts sent:	16089	ack pkts sent:	30790
pure acks sent:	16086	pure acks sent:	4
sack pkts sent:	1466	sack pkts sent:	0
dsack pkts sent:	0	dsack pkts sent:	0
max sack blks/ack:	2	max sack blks/ack:	0
unique bytes sent:	1515	unique bytes sent:	38719841
actual data pkts:	2	actual data pkts:	30785
actual data bytes:	1515	actual data bytes:	38785361
rexmt data pkts:	0	rexmt data pkts:	52
rexmt data bytes:	0	rexmt data bytes:	65520
zwnd probe pkts:	0	zwnd probe pkts:	0
zwnd probe bytes:	0	zwnd probe bytes:	0
outoforder pkts:	0	outoforder pkts:	16
pushed data pkts:	1	pushed data pkts:	578
SYN/FIN pkts sent:	1/1	SYN/FIN pkts sent:	1/1
req 1323 ws/ts:	Y/N	req 1323 ws/ts:	Y/N
adv wind scale:	1	adv wind scale:	6
req sack:	Y	req sack:	Y
sacks sent:	1466	sacks sent:	0
urgent data pkts:	0 pkts	urgent data pkts:	0 pkts
urgent data bytes:	0 bytes	urgent data bytes:	0 bytes
mss requested:	1260 bytes	mss requested:	1380 bytes
max segm size:	1260 bytes	max segm size:	1260 bytes
min segm size:	255 bytes	min segm size:	16 bytes
avg segm size:	757 bytes	avg segm size:	1259 bytes
max win adv:	65536 bytes	max win adv:	8896 bytes
min win adv:	16 bytes	min win adv:	8896 bytes
zero win adv:	3 times	zero win adv:	0 times
avg win adv:	65251 bytes	avg win adv:	8896 bytes
initial window:	1515 bytes	initial window:	2197 bytes
initial window:	2 pkts	initial window:	3 pkts
ttl stream length:	1515 bytes	ttl stream length:	38719841 bytes
missed data:	0 bytes	missed data:	0 bytes
truncated data:	1443 bytes	truncated data:	37677121 bytes
truncated packets:	2 pkts	truncated packets:	30784 pkts

data xmit time:	0.000 secs	data xmit time:	111.603 secs
idletime max:	708.8 ms	idletime max:	708.9 ms
throughput:	14 Bps	throughput:	345248 Bps
RTT samples:	3	RTT samples:	14596
RTT min:	30.6 ms	RTT min:	49.1 ms
RTT max:	31.1 ms	RTT max:	212.6 ms
RTT avg:	30.8 ms	RTT avg:	49.9 ms
RTT stdev:	0.3 ms	RTT stdev:	3.6 ms
RTT from 3WHS:	31.1 ms	RTT from 3WHS:	49.8 ms
RTT full_sz smpls:	2	RTT full_sz smpls:	1
RTT full_sz min:	30.6 ms	RTT full_sz min:	49.8 ms
RTT full_sz max:	31.1 ms	RTT full_sz max:	49.8 ms
RTT full_sz avg:	30.9 ms	RTT full_sz avg:	49.8 ms
RTT full_sz stdev:	0.0 ms	RTT full_sz stdev:	0.0 ms
post-loss acks:	0	post-loss acks:	52
segs cum acked:	1	segs cum acked:	16086
duplicate acks:	2	duplicate acks:	1435
triple dupacks:	0	triple dupacks:	52
max # retrans:	0	max # retrans:	1
min retr time:	0.0 ms	min retr time:	79.9 ms
max retr time:	0.0 ms	max retr time:	158.0 ms
avg retr time:	0.0 ms	avg retr time:	88.1 ms
sdv retr time:	0.0 ms	sdv retr time:	22.5 ms

Throughput : b->a

Downlink Throughput= (unique bytes sent * 8)/1,000,000

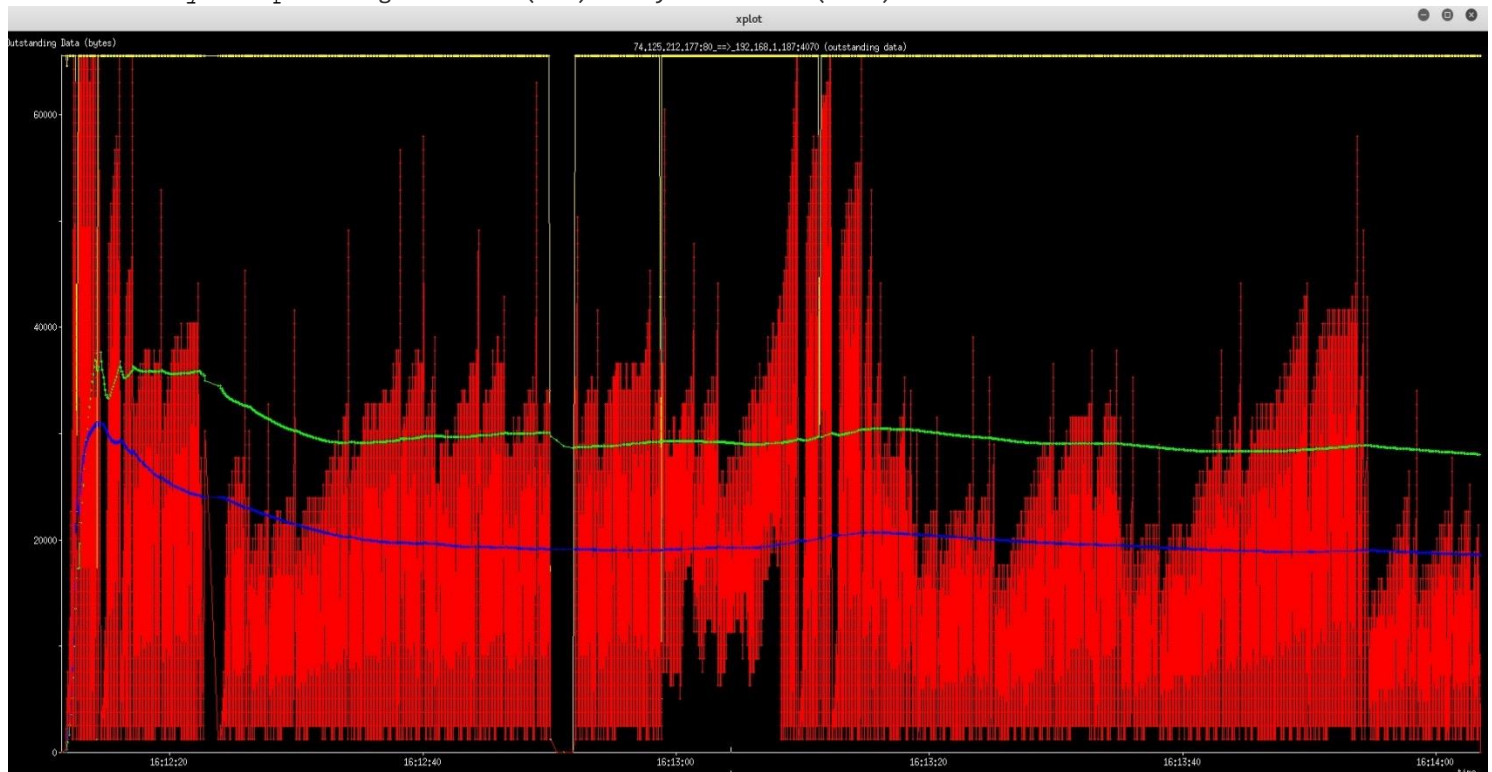
= 2.775Mbps (According to tcptrace, Wireshark reports slightly different value)

Congestion Window: (Explanation same as Wireshark)


Command:

xplot.org b2a_owin.xpl (plots unacked bytes for b->a)

Un-ACKed Bytes plot: green line (ACK) and yellow line (RWIN)



Time Sequence Plot:



2c.

Zero Window

Command:

```
tcptrace -r -l pa2_trace.pcap
```

The output of the above command gives the Zero Window Advertisement(if any).The output is given in part 2b and we can observe that the Local Host (192.168.1.187), makes a Zero Window Adv 3 times.

From Receivers point of view (192.168.1.187,4070,74.125.212.177,80)

Iperf

- Measure the bi-directional bandwidth between client and server at 3 different locations

Location 1

```
G:\Users\joo5819\Downloads\iperf-3.0.11-win64>iperf3.exe -c 192.168.1.83 -d
send_parameters:
{
    "tcp": true,
    "omit": 0,
    "time": 10,
    "parallel": 1,
    "len": 131072,
    "client_version": "3.1b3"
}
Connecting to host 192.168.1.83, port 5201
SO_SNDBUF is 212992
[ 4] local 192.168.1.90 port 49226 connected to 192.168.1.83 port 5201
[ ID] Interval           Transfer     Bandwidth
[ 4]   0.00-1.00      sec    2.75 MBytes    22.9 Mbits/sec
[ 4]   1.01-2.00      sec    2.00 MBytes    16.9 Mbits/sec
[ 4]   2.00-3.01      sec    2.38 MBytes    19.8 Mbits/sec
[ 4]   3.01-4.00      sec    2.75 MBytes    23.2 Mbits/sec
[ 4]   4.00-5.01      sec    2.88 MBytes    24.0 Mbits/sec
[ 4]   5.01-6.01      sec    2.50 MBytes    20.8 Mbits/sec
[ 4]   6.01-7.01      sec    2.62 MBytes    22.2 Mbits/sec
[ 4]   7.01-8.00      sec    2.50 MBytes    21.1 Mbits/sec
[ 4]   8.00-9.00      sec    2.62 MBytes    22.0 Mbits/sec
send_results
{
    "cpu_util_total": 22.3658,
    "cpu_util_user": 0.455258,
    "cpu_util_system": 21.9105,
    "sender_has_retransmits": 0,
    "streams": {
        "id": 1,
        "bytes": 26738688,
        "retransmits": -1
    }
}
```

```
    "streams": {
        "id": 1,
        "bytes": 26738688,
        "retransmits": -1,
        "jitter": 0,
        "errors": 0,
        "packets": 0
    }
}
[ 4]   9.00-10.00    sec    2.50 MBytes    21.0 Mbits/sec
- - - - -
[ ID] Interval           Transfer     Bandwidth
[ 4]   0.00-10.00     sec    25.5 MBytes    21.4 Mbits/sec
[ 4]   0.00-10.00     sec    25.5 MBytes    21.4 Mbits/sec
iperf Done.
```

Location 2

```
{
  "cpu_util_total": 3.74733,
  "cpu_util_user": 0.59573,
  "cpu_util_system": 3.1516,
  "sender_has_retransmits": 0,
  "streams": [
    {
      "id": 1,
      "bytes": 19136512,
      "retransmits": -1,
      "jitter": 0,
      "errors": 0,
      "packets": 0
    }
  ]
}
get_results
{
  "cpu_util_total": 0.768734,
  "cpu_util_user": 0.106403,
  "cpu_util_system": 0.662319,
  "sender_has_retransmits": -1,
  "streams": [
    {
      "id": 1,
      "bytes": 19136512,
      "retransmits": -1,
      "jitter": 0,
      "errors": 0,
      "packets": 0
    }
  ]
}
[ 4] 9.00-10.00 sec 1.88 MBytes 15.7 Mbits/sec
-----
[ ID] Interval      Transfer      Bandwidth
[ 4] 0.00-10.00 sec 18.2 MBytes 15.3 Mbits/sec      sender
[ 4] 0.00-10.00 sec 18.2 MBytes 15.3 Mbits/sec      receiver

iperf Done.
C:\Users\yakshdeep\Desktop\iperf-3.0.11-win64>
```

Location 3

```
C:\Users\yakshdeep\Desktop\iperf-3.0.11-win64>
{
  "tcp": true,
  "omit": 0,
  "time": 10,
  "parallel": 1,
  "len": 131072,
  "client_version": "3.1b3"
}
>
Connecting to host 192.168.1.85, port 5201
SO_SNDBUF is 212992
[ 4] local 192.168.1.90 port 56965 connected to 192.168.1.85 port 5201
[ ID] Interval      Transfer      Bandwidth
[ 4] 0.00-1.00 sec 3.75 MBytes 31.4 Mbits/sec
[ 4] 1.00-2.02 sec 4.12 MBytes 34.1 Mbits/sec
[ 4] 2.02-3.01 sec 4.50 MBytes 38.1 Mbits/sec
[ 4] 3.01-4.00 sec 4.25 MBytes 35.8 Mbits/sec
[ 4] 4.00-5.00 sec 4.50 MBytes 37.8 Mbits/sec
[ 4] 5.00-6.00 sec 4.62 MBytes 38.6 Mbits/sec
[ 4] 6.00-7.00 sec 4.00 MBytes 33.7 Mbits/sec
[ 4] 7.00-8.01 sec 4.75 MBytes 39.5 Mbits/sec
[ 4] 8.01-9.02 sec 4.25 MBytes 35.2 Mbits/sec
send_results
{
  "cpu_util_total": 41.7309,
  "cpu_util_user": 0.761585,
  "cpu_util_system": 40.9694,
  "sender_has_retransmits": 0,
  "streams": [
    {
      "id": 1,
      "bytes": 44957696,
      "retransmits": -1,
      "jitter": 0,
      "errors": 0,
      "packets": 0
    }
  ]
}
>1
get_results
{
  "cpu_util_total": 1.14028,
  "cpu_util_user": 0.129477,
  "cpu_util_system": 1.01081,
  "sender_has_retransmits": -1,
  "streams": [
    {
      "id": 1,
      "bytes": 44955104,
      "retransmits": -1,
      "jitter": 0,
      "errors": 0,
      "packets": 0
    }
  ]
}
>1
[ 4] 9.02-10.00 sec 4.12 MBytes 35.2 Mbits/sec
-----
[ ID] Interval      Transfer      Bandwidth
[ 4] 0.00-10.00 sec 42.9 MBytes 36.0 Mbits/sec      sender
[ 4] 0.00-10.00 sec 42.9 MBytes 35.9 Mbits/sec      receiver
```

- Change TCP window size appropriately to get better bandwidth.

Window Size: 5000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
<i>11.5</i>	<i>9.61</i>
<i>13.9</i>	<i>11.7</i>
<i>12.1</i>	<i>10.1</i>
<i>13.9</i>	<i>11.6</i>

Window Size 10000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
<i>16.8</i>	<i>14.1</i>
<i>15.3</i>	<i>12.8</i>
<i>11.8</i>	<i>9.9</i>
<i>15.1</i>	<i>12.7</i>

Widow Size 16357 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
<i>18.3</i>	<i>15.3</i>
<i>19.2</i>	<i>16.1</i>
<i>14.8</i>	<i>12.4</i>
<i>16.9</i>	<i>14.2</i>

Window Size 25000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
<i>14.7</i>	<i>12.3</i>
<i>17.2</i>	<i>14.4</i>
<i>15.7</i>	<i>13.2</i>
<i>18.9</i>	<i>15.8</i>

Window Size 40000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
23.7	19.9
21.7	18.2
22.5	18.8
24.5	20.5

Window Size 50000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
20.4	17.1
22.9	19.2
19.6	16.4
24.2	20.3

Window Size 60000 bytes:

<i>Transfer Data Size(Mb)</i>	<i>Bandwidth(Mb/s)</i>
23.6	19.8
25.9	21.7
22.1	18.6
22.8	19.2

Explanation

For best Results, tcp window size = Bandwidth delay product. In our case, the bandwidth that we had was 34Mbps & average delay of 10ms, so BDP= 42500 Bytes.

Thus, we can see from the above results that the bandwidth keeps improving till 40KB (window size) and remains somewhat constant for window sizes >40KB.

- Compare the throughput observed when using bi-directional UDP and a TCP

UDP Result

```
C:\Users\yakshdeep\Desktop\iperf-3.0.11-win64>iperf3.exe -c 24.72.251.234 -u -b 10000m
Connecting to host 24.72.251.234, port 5201
[ 4] local 24.72.250.122 port 58139 connected to 24.72.251.234 port 5201
[ ID] Interval           Transfer     Bandwidth       Total Datagrams
[ 4]  0.00-1.02   sec    1.88 MBytes   15.5 Mbits/sec    241
[ 4]  1.02-2.02   sec    1.98 MBytes   16.6 Mbits/sec    253
[ 4]  2.02-3.00   sec    2.15 MBytes   18.3 Mbits/sec    275
[ 4]  3.00-4.01   sec    2.23 MBytes   18.6 Mbits/sec    285
[ 4]  4.01-5.00   sec    2.22 MBytes   18.7 Mbits/sec    284
[ 4]  5.00-6.01   sec    2.02 MBytes   16.8 Mbits/sec    259
[ 4]  6.01-7.01   sec    1.93 MBytes   16.3 Mbits/sec    247
[ 4]  7.01-8.01   sec    1.72 MBytes   14.4 Mbits/sec    220
[ 4]  8.01-9.00   sec    1.72 MBytes   14.5 Mbits/sec    220
[ 4]  9.00-10.00  sec    2.05 MBytes   17.2 Mbits/sec    262
-----
[ ID] Interval           Transfer     Bandwidth       Jitter    Lost/Total Datagrams
[ 4]  0.00-10.00  sec   19.9 MBytes   16.7 Mbits/sec   3.479 ms  1/2544 (0.039%)
[ 4] Sent 2544 datagrams

iperf Done.
```

TCP Result:

```
send_results
{
    "cpu_util_total":      1.21578,
    "cpu_util_user":       0,
    "cpu_util_system":     1.21578,
    "sender_has_retransmits": 0,
    "streams":             [{
        "id": 1,
        "bytes": 14362144,
        "retransmits": -1,
        "jitter": 0,
        "errors": 0,
        "packets": 0
    }]
}
get_results
{
    "cpu_util_total":      1.97622,
    "cpu_util_user":       0.69041,
    "cpu_util_system":     1.28581,
    "sender_has_retransmits": -1,
    "streams":             [{
        "id": 1,
        "bytes": 14316164,
        "retransmits": -1,
        "jitter": 0,
        "errors": 0,
        "packets": 0
    }]
}
[ 4]  9.00-10.00  sec    1.72 MBytes   14.4 Mbits/sec
-----
[ ID] Interval           Transfer     Bandwidth
[ 4]  0.00-10.00  sec   13.7 MBytes   11.5 Mbits/sec      sender
[ 4]  0.00-10.00  sec   13.7 MBytes   11.5 Mbits/sec      receiver
```

UDP performs better than TCP in the tested network conditions since in UDP mode, iperf has no flow control whereas TCP does (limiting bandwidth). Since the data rate is not high in this case, UDP performs better than TCP.

For higher data rates, it is imperative that TCP performs better.

- Install iperf in your phone(server) . Analyze the bandwidth measured when the server is moving.

```

Interface: wlan0 ipaddr: 143.215.105.36
SGH-M919 (MSM8960) [ARM]
WLAN: UnknownSGH-M919 (MSM8960) [ARM]
WLAN: UnknownSGH-M919 (MSM8960) [ARM]
WLAN: UnknownSGH-M919 (MSM8960) [ARM]

ipref - s
/data/data/com.magicandroidapps.iperf/bin/iperf:
ignoring extra argument -- ipref

-----
Server listening on TCP port 5001
TCP window size: 1.00 MByte (default)
-----
[ 4] local 143.215.105.36 port 5001 connected with
143.215.100.110 port 63449
[ ID] Interval Transfer Bandwidth[ 5] local
143.215.105.36 port 5001 connected with
143.215.100.110 port 63491[ 5] 0.0-37.8 sec 14.0 Bytes
2.96 bits/sec[ 4] local 143.215.105.36 port 5001
connected with 143.215.100.110 port 63492[ 4]
0.0-200.0 sec 14.0 Bytes 0.56 bits/sec/data/data/

```

The speeds are terrible when the server is mobile this is probably due to the server getting farther away from the access point.